

Computing Arbitrary Functions of Encrypted Data

Klev

September 14, 2015

- Fully homomorphic encryption scheme.
 - Keep data private, only people with access can decrypt it.
 - Still, possible to compute on the encrypted data.
 - Cloud computing compatible with privacy
- Fully homomorphic encryption
 - Fully: No restrictions on the operations that can be performed
 - Send description of function f to the server. Server computes the function on the encrypted data. Returns encrypted data. User decrypts.
 - Useful whenever the response can be encrypted.
 - Alice and her jewelry store.

1 Homomorphic Encryption: Functionality

Three algorithms: *KeyGen*, *Encrypt*, *Decrypt*

Symmetric encryption scheme: Key used for both encryption and decryption

- Asymmetric encryption scheme: Public key for encryption, Private key for decryption

- HE can be either symmetric or asymmetric. Focus on asymmetric.

- Fourth algorithm: *Evaluate*. A set of permitted functions F_e . *Evaluate* can handle functions in F_e . Others are undefined.

2 Requirements

- Require that decrypting c (output from *Evaluate*) takes the same amount of time as decrypting c_1 (output from *Encrypt*).
 - c is the same size as c_1
 - "Compact ciphertext requirements"

- Size of c and the time required to decrypt it is independent of f
- Fully homomorphic if it can handle all functions, has compact ciphertexts and *Evaluate* is efficient.
- The work required by Alice to extract jewelry has no connection to the time spent by workers to assemble the piece.
- Measure complexity of function f : Running time T_f and size S_f of a boolean circuit.
- Break down the computation of f into AND, OR and NOT gates.
- Obtain fully homomorphic encryption by operating on ciphertexts using add, subtract and multiply.
- If encrypted, random-access can not speed up the work. *Evaluate* must touch all points. Runtime is at least linear.
- Not touching them would leak information, saying that they are irrelevant for function f .
- Trade off: the data is contained in 1% of my area. Reduce computation by factor 100.
- Should be able to specify the amount of data received from query. Unavoidable with padding and truncating.

3 Homomorphic Encryption: Security

- Semantic security against chosen-plaintext attacks (CPA)
 - Given c and two messages m_1 and m_2 , it is *hard* for an adversary to determine which plaintext the ciphertext decrypts to.
 - "Hard": Guesses correctly with probability $1/2 + e = \text{Advantage}$
 - Deterministic: Same plaintext encrypts to same ciphertext every time.
 - The scheme must be probabilistic; a plaintext encrypts to multiple ciphertexts.
 - Malleability weakens deterministic schemes, not semantically secure.

4 A Somewhat Homomorphic Encryption Scheme

- First as a symmetric encryption scheme.
 - $N = \lambda$, $P = \lambda^2$, $Q = \lambda^5$
 - KeyGen: Key is random P -bit odd integer p
 - Encrypt(e, m): $m' = \text{random } N\text{-bit number such that } m' = m \bmod 2$.
 $c \leftarrow m' + pq$
 - Decrypt(p, c): $(cm \bmod p) \bmod 2$, where $cm \bmod p$ is the integer c' in $(-p/2, p/2)$ such that p divides $c - c'$.
 - $cm \bmod p$ is the noise associated to a ciphertext. Distance to nearest multiple of p .
 - Consider $Mult_e$: $c = c_1 * c_2$. Noise to $c_i = m'_i$. We have that $c = m'_1 * m'_2 + pq'$

- Multiplication tends to increase the noise faster than addition and subtraction.
- c_1 and c_2 with k_1 and k_2 -bit noises = roughly $(k_1 + k_2)$ -bit noises.
- The functions that can be handled are those where $|f^\dagger(a_1, \dots, a_t)|$ is always less than $p/2$.

5 Bootstrappable Encryption

- Work on box 1, put it inside box 2 which contains key 1, open box 1 and continue work through box 2.
- Somewhat homomorphic encryption: Supports Add, Subtract, Multiply, and a limited set of operations, until the noise gets too large.
- Key for box i represents an *encrypted secret decryption key*.
- Most important function: The decryption function (Open up the inner box to continue work)
- Decrypts itself \rightarrow bootstrappable

5.1 Bootstrappable to Fully Homomorphic

- Suppose e can handle decryption, ADD, SUB and MULT.
- If these four algorithms are in f_e , we can construct an encryption scheme that is fully homomorphic.
- $Recrypt_e(pk_2, D_e, sk_1, c_1)$ Recrypt outputs an encryption of m , but under the new key pk_2 . The message is encrypted twice under different keys pk_1 and pk_2 .
- Can "peel off" the layers like an onion, but the *evaluate* algorithm "opens" the inner encryption. (Like in the box example).
- Decryption the inner layer removes noise. Using *Evaluate* with pk_2 adds new noise. If the added noise is less than what we remove in the decryption, we've made progress.
- Public key consists of a sequence of public keys $(pk_1, pk_2, \dots, pk_n)$ encrypted under secret keys sk_{i+1}

5.2 Circular Security

- Safe to reveal the encryption of a secret key s_k under its own associated public key p_k .

5.3 Greasing the Decryption Circuit

- $m \leftarrow (c \bmod p) \bmod 2 == m \leftarrow LSB(c) \oplus LSB(\lfloor c/p \rfloor)$
- If not $\lfloor c/p \rfloor$ is complicated, then e is bootstrappable and FHE is possible. Not possible with the current scheme. Let's create e^*
- Replace e 's decryption function, which multiplies two long numbers, with a decryption function that adds a fairly small set of numbers.

5.3.1 Transformation

- KeyGen: Run the algorithm to obtain p_k and s_k (s_k is an odd integer p).
 - Generate a subset from s_k of rational numbers that sums up to roughly $1/p$. s_k will then be the subset y of the numbers encoded as a vector. $p_k* \leftarrow (p_k, y)$
 - This is a "hint" about the secret key. Adds the "grease" to the system. Revealing this obviously impacts security.
- Encrypt: Run $\text{Encrypt}(p_k, m)$ to obtain ciphertext c . For $i \in \{1, \dots, \beta\}$, set $z_i \leftarrow c * y_i \text{ mod } 2$. The ciphertext $c*$ is now c and $z = \{z_1, \dots, z_\beta\}$
 - The hint is used to postprocess the ciphertext. The idea is to leave less work for the Decrypt algorithm.
 - Used in 'Server-aided cryptography'.
 - The hint is statistically dependent on s_k , but enough for the server to decrypt efficiently on its own.
 - In our case: The encrypter or evaluator plays the role as server
- Decrypt: Output $LSB(c) \oplus LSB(\lfloor \sum s_i z_i \rfloor)$
 - Works because $\sum s_i z_i = \sum c * s_i y_i = c/p \text{ mod } 2$
 - Important: Replace the multiplication of c and $1/p$ with a summation that contains only α nonzero terms.
- Add: (p_k*, c_1*, c_2*) . Extract c_1 and c_2 from c_1* and c_2* . Run $c \leftarrow \text{Add}(p_k, c_1, c_2)$. The result consist of ciphertext $c*$. $c* = c$ and the result of postprocessing c with y .

5.4 How to Add Numbers

- Let us consider the computation of $\sum s_i z_i$.
 -

5.5 Security of the Transformed Scheme

- The encryption key contains a hint about the secret p .